

The Introduction to FEMU and its newly related research

Presented by Chen Haodong

Outline

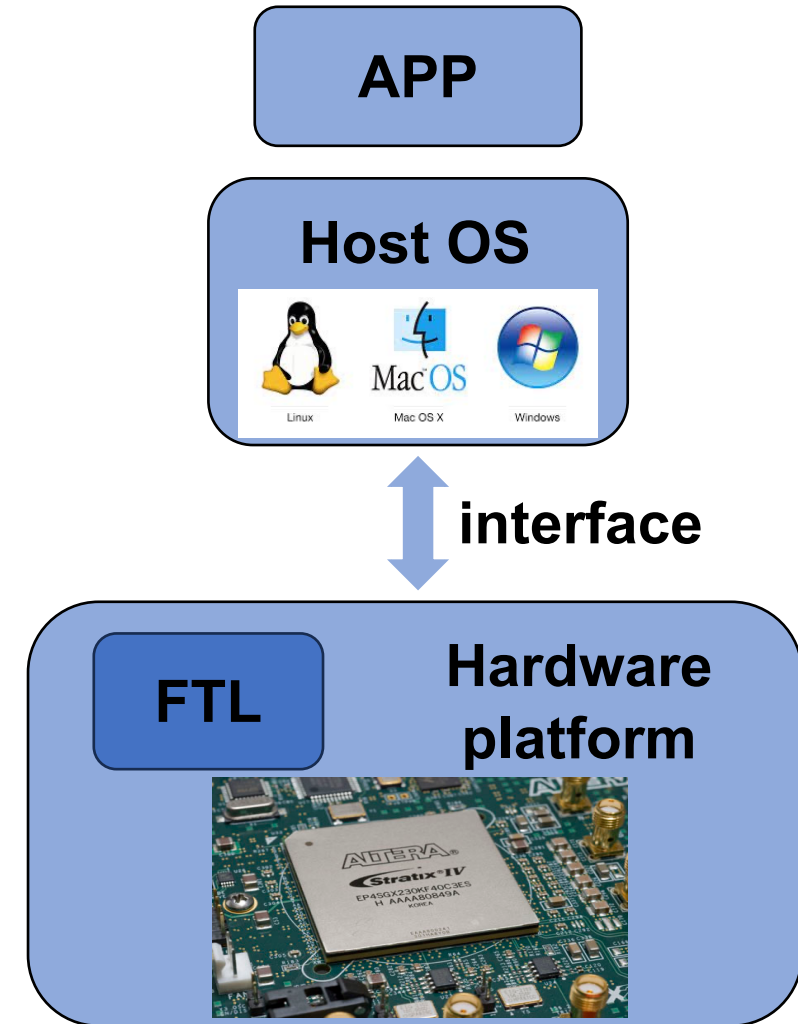
- **1. Backgrounds**
- **2. Existing SSD platform**
- **3. Case of FEMU**
- **4. Newly related work in 2023**
- **5. Future work**

1. Backgrounds

SSD research mainly focus on 3 parts:

1. kernel changes
2. Interface changes
3. FTL changes

Typical full-stack research



HDU – Center of Cloud computing

2. Existing SSD platforms

Simulator

Emulator

**Hardware
platform**

2. Existing SSD platforms

Simulator

DiskSim's SSD;
FlashSim; SSDSim;

Emulator

Hardware platform

2. Existing SSD platforms

Simulator

DiskSim's SSD;
FlashSim; SSDSim;

Emulator

Hardware platform



simple



Time-saving



**internal-SSD
research**

2. Existing SSD platforms

Simulator

DiskSim's SSD;
FlashSim; SSDSim;

Emulator

Hardware platform



simple



Time-saving



**internal-SSD
research**



**kernel-level
research**

2. Existing SSD platforms

Simulator

DiskSim's SSD;
FlashSim; SSDSim;

Emulator

Hardware platform

OpenSSD; FPGA;
OpenChannel SSD;

-  simple
-  Time-saving
-  internal-SSD research
-  kernel-level extension

2. Existing SSD platforms

Simulator



DiskSim's SSD;
FlashSim; SSDSim;

Emulator

Hardware platform

OpenSSD; FPGA;
OpenChannel SSD;

-  simple
-  Time-saving
-  internal-SSD research
-  kernel-level extension

-  Full-stack research
-  Accurate

2. Existing SSD platforms


Simulator

DiskSim's SSD;
FlashSim; SSDSim;

 simple

 Time-saving

 internal-SSD
research

 kernel-level
extension

Emulator

Hardware platform

OpenSSD; FPGA;
OpenChannel SSD;

 Full-stack
research

 Accurate

 high costs

 complex to use

2. Existing SSD platforms

Simulator

DiskSim's SSD;
FlashSim; SSDSim;

Emulator

FlashEm;
LightNVM QEMU;
QEMU based VSSIM;

Hardware platform

OpenSSD; FPGA;
OpenChannel SSD;



simple



Time-saving



**internal-SSD
research**



**kernel-level
extension**



**Full-stack
research**



Accurate



high costs



complex to use

2. Existing SSD platforms


Simulator

DiskSim's SSD;
FlashSim; SSDSim;

 simple

 Time-saving

 internal-SSD
research

 kernel-level
extension

Emulator

FlashEm;
LightNVM QEMU;
QEMU based VSSIM;

 Full-stack
research

 Low costs

Hardware platform

OpenSSD; FPGA;
OpenChannel SSD;

 Full-stack
research

 Accurate

 high costs

 complex to use

2. Existing SSD platforms


Simulator

DiskSim's SSD;
FlashSim; SSDSim;

 simple

 Time-saving

 internal-SSD
research

 kernel-level
extension

Emulator

FlashEm;
LightNVM QEMU;
QEMU based VSSIM;

 Full-stack
research

 Low costs

 Not scalability

 Not accurate

Hardware platform

OpenSSD; FPGA;
OpenChannel SSD;

 Full-stack
research

 Accurate

 high costs

 complex to use

3. The CASE of FEMU

FEMU is a software (QEMU-based) flash emulator for fostering future full-stack software/hardware SSD research.

3. The *CASE* of *FEMU*

FEMU is a software (QEMU-based) flash emulator for fostering future full-stack software/hardware SSD research.



Cheap

**Open source,
download from:**

<https://github.com/vtess/FEMU>

3. The CASE of FEMU

FEMU is a software (QEMU-based) flash emulator for fostering future full-stack software/hardware SSD research.



Cheap



**(relatively)
Accurate**

**Produce an advanced openChannel SSD
Model**

**The error of resulting latencies drops to
only 0.5-38%**

3. The *CASE* of *FEMU*

FEMU is a software (QEMU-based) flash emulator for fostering future full-stack software/hardware SSD research.



Cheap



**(relatively)
Accurate**



Scalable

support 32 channels/chips

3. The *CASE* of *FEMU*

FEMU is a software (QEMU-based) flash emulator for fostering future full-stack software/hardware SSD research.



Cheap



**(relatively)
Accurate**



Scalable



Extensible

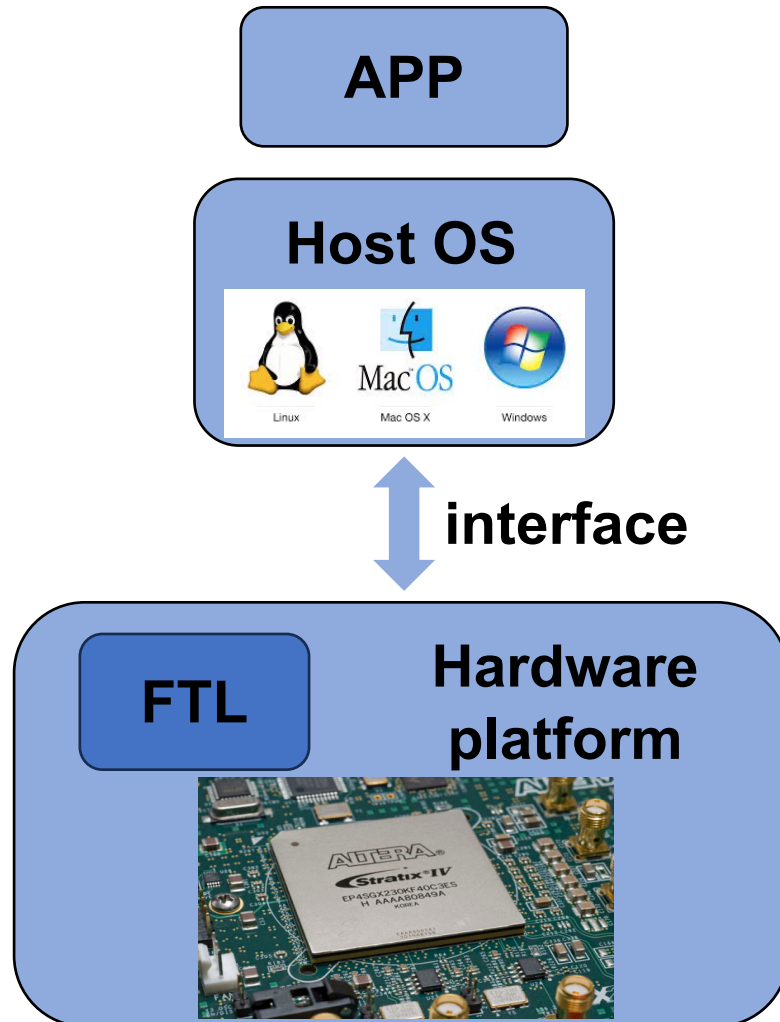
- **modifiable interface**
 - **modifiable FTL**
 - **modifiable SSD**
- **modifiable kernel**

3. The CASE of FEMU

Typical full stack research

.vs


FEMU full stack research



HDU – Center of Cloud computing

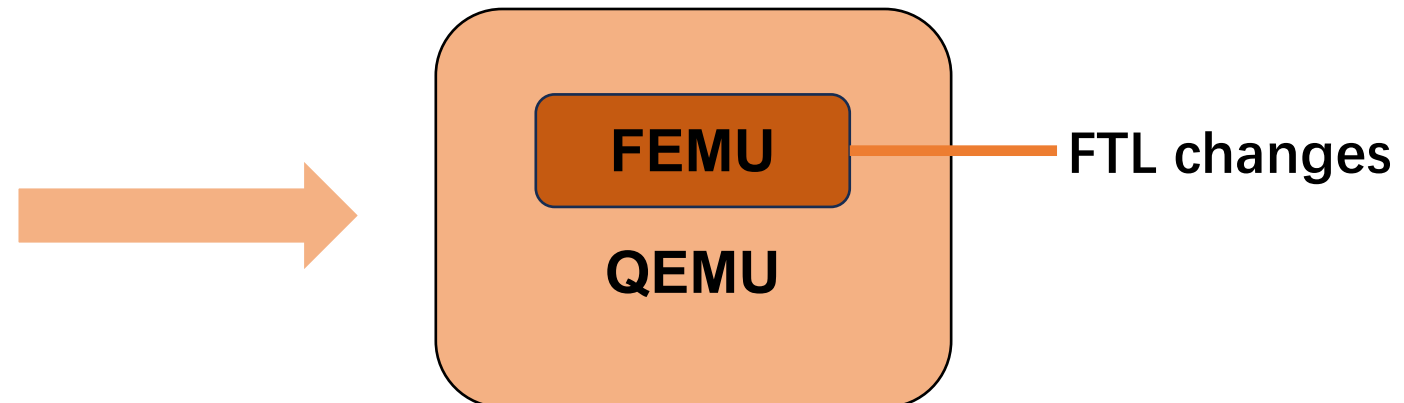
3. The CASE of FEMU

Typical full stack research



.vs

FEMU full stack research


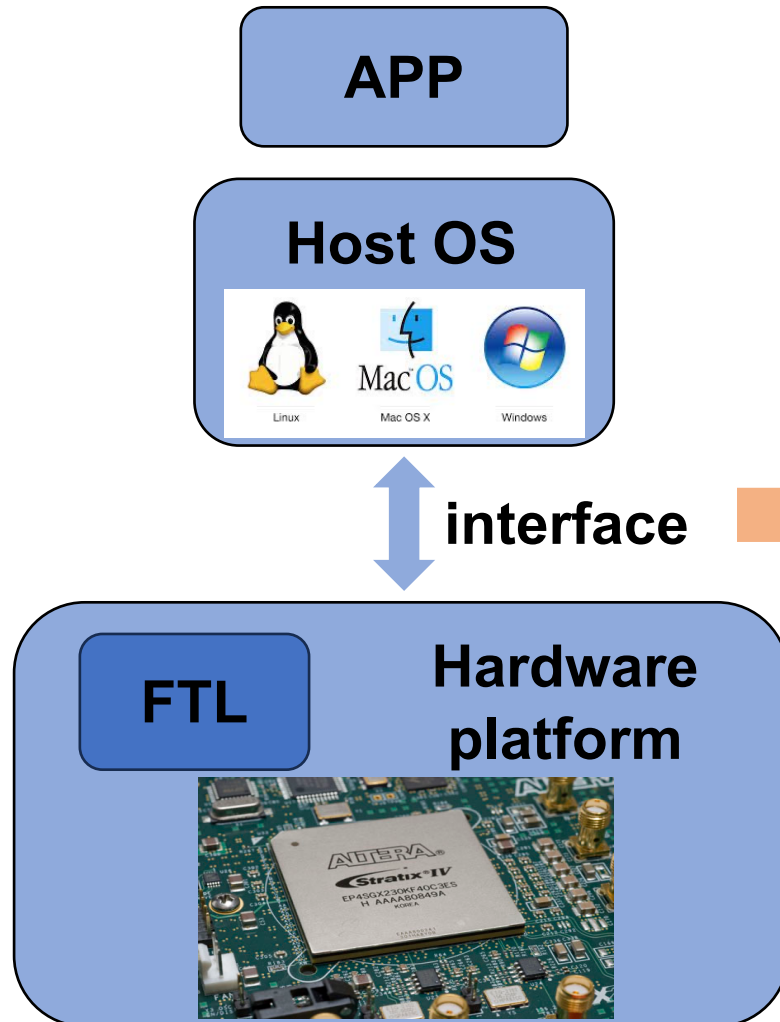


3. The CASE of FEMU

Typical full stack research

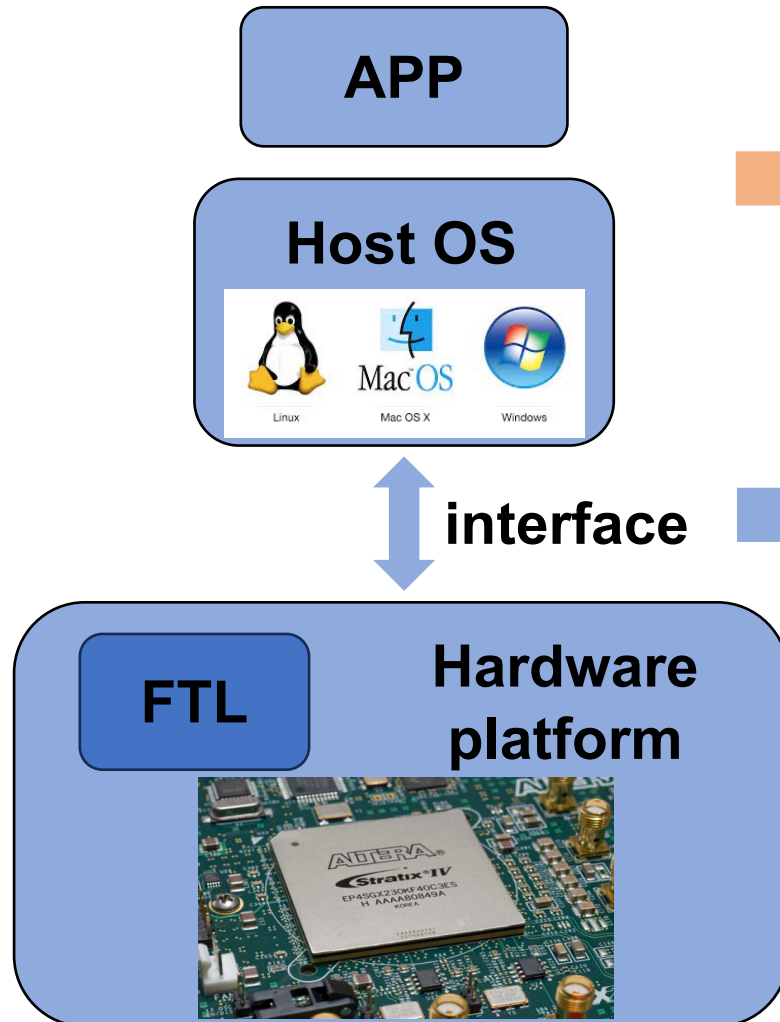
.vs

FEMU full stack research




3. The CASE of FEMU

Typical full stack research



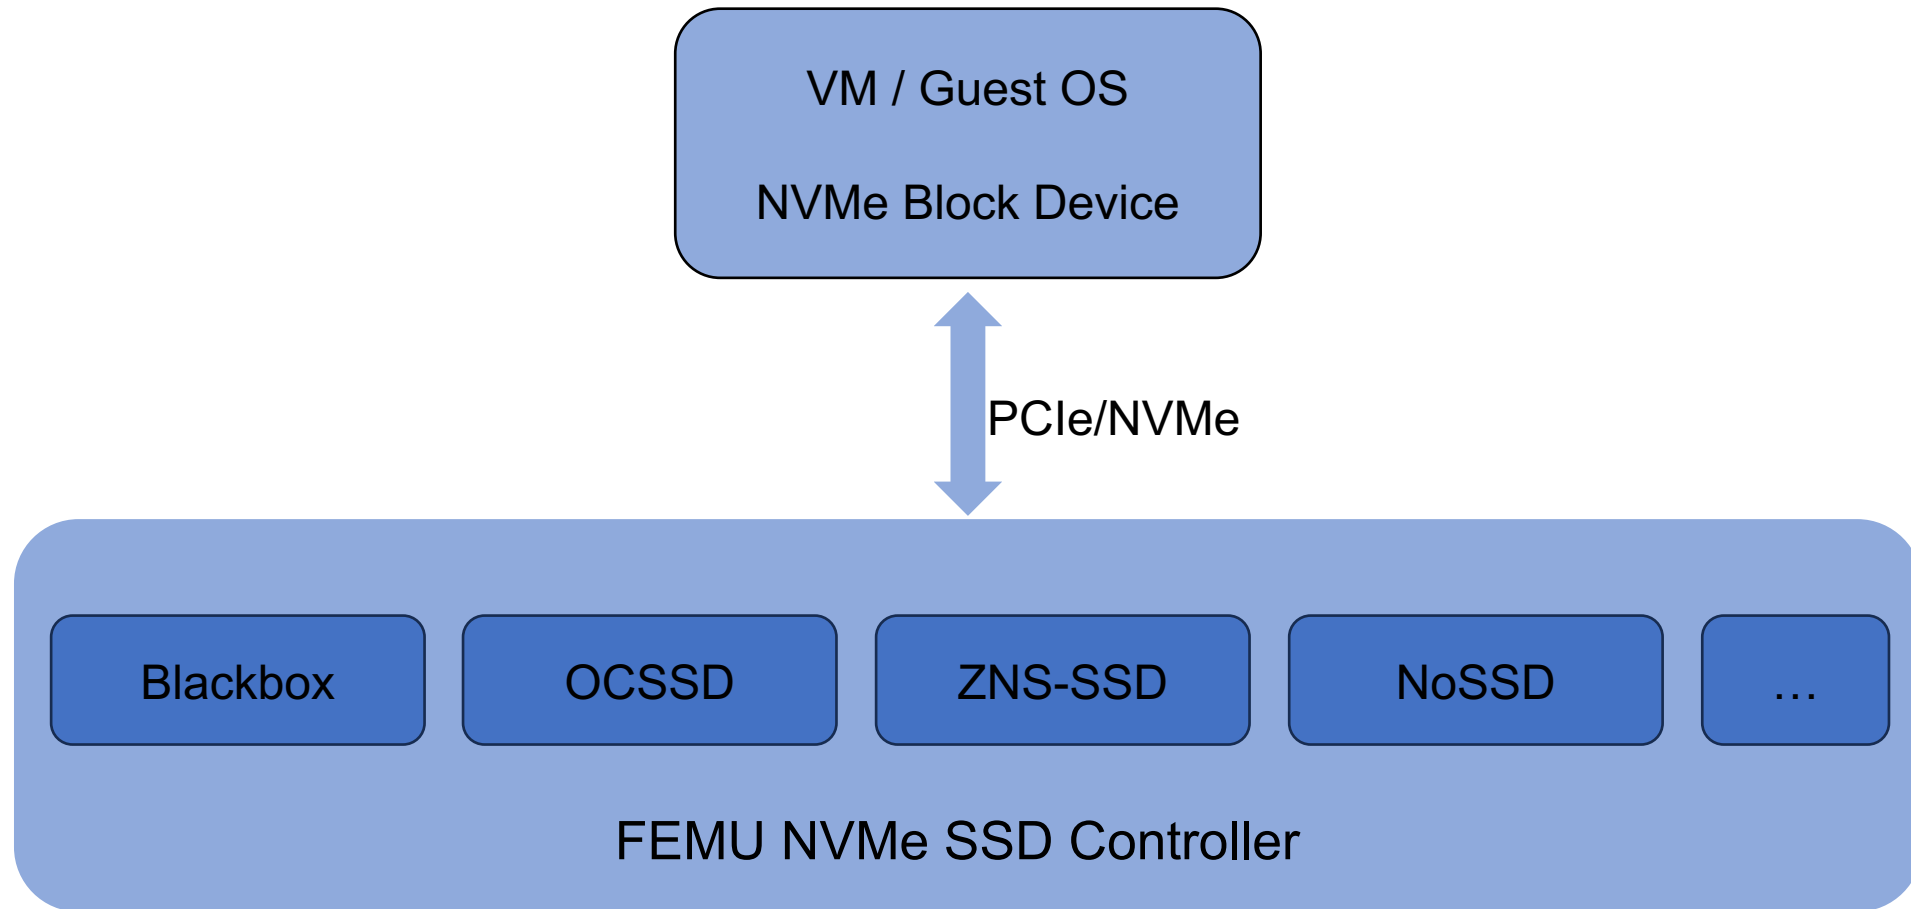
.vs

FEMU full stack research



3. The CASE of FEMU

Project Description



4. Newly related work in 2023

1. **Holistic and Opportunistic Scheduling of Background I/Os in Flash-based SSDs. (TC'23).** Yu Wang, You Zhou, Fei Wu, Yu Zhong, Jian Zhou, Zhonghai Lu, Shu Li, Zhenghong Wang, Changsheng Xie

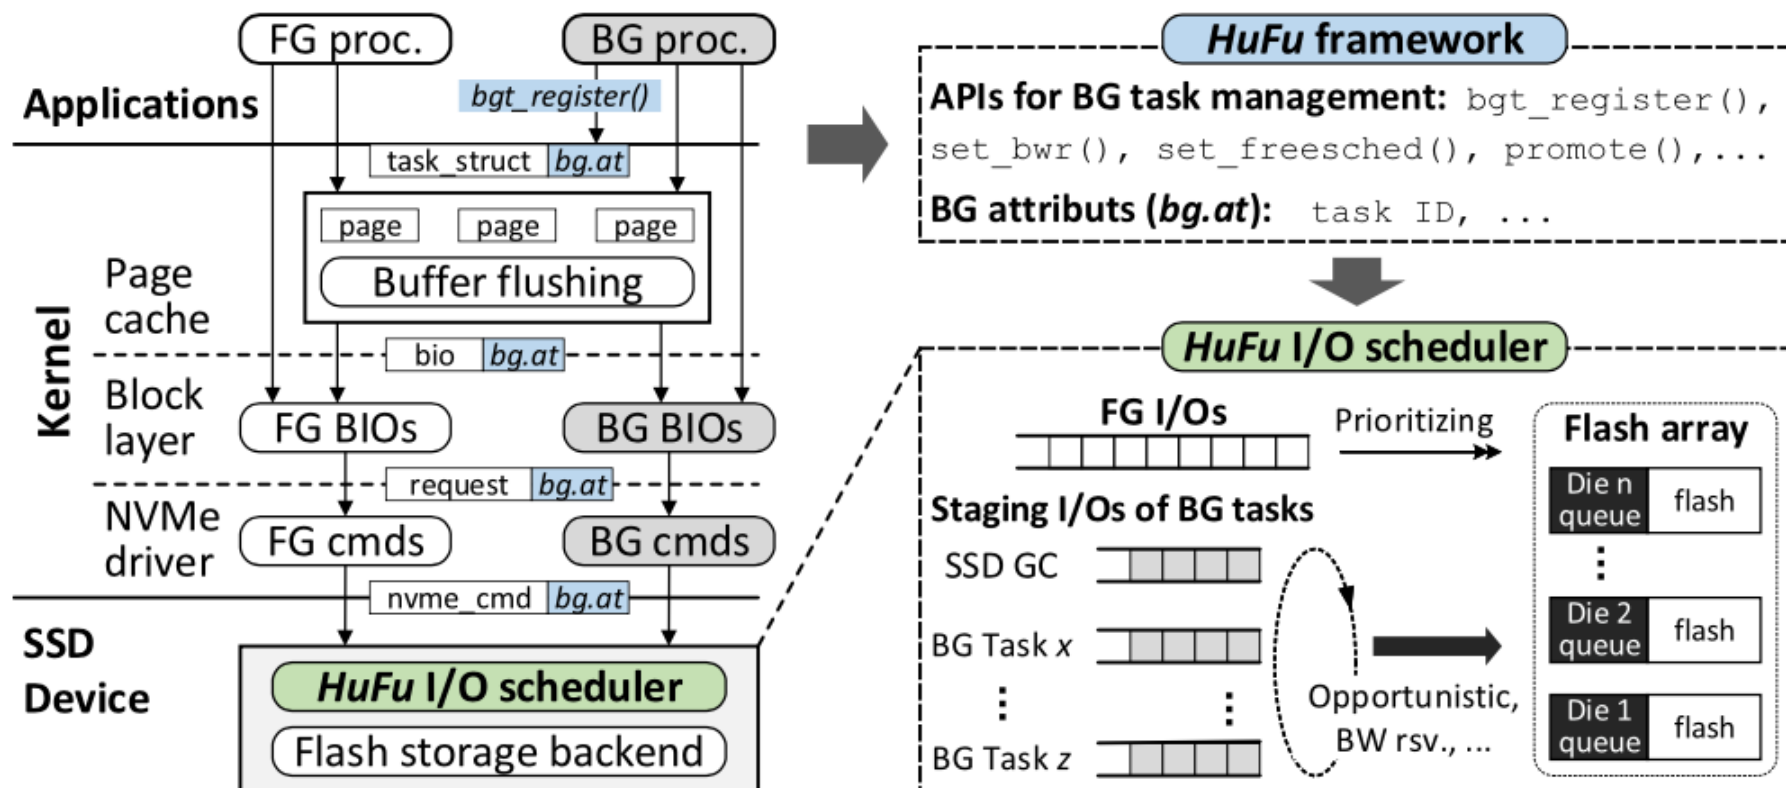



Fig. 3. Design overview of *HuFu*. The core components include a framework, which allows host software to offload I/O scheduling of BG tasks into the SSD, and an SSD-internal I/O scheduler, which prioritizes FG I/Os while processing system-wide BG I/Os opportunistically.

4. Newly related work in 2023

1. **Holistic and Opportunistic Scheduling of Background I/Os in Flash-based SSDs. (TC'23).** Yu Wang, You Zhou, Fei Wu, Yu Zhong, Jian Zhou, Zhonghai Lu, Shu Li, Zhenghong Wang, Changsheng Xie



FEMU

Kernel changes

FTL changes

HDU – Center of Cloud computing

Fig. 3. Design overview of *HuFu*. The core components include a framework and an SSD-internal I/O scheduler, which prioritizes FG I/Os while processing s

4. Newly related work in 2023

2. **NVMeVirt: A Versatile Software-defined Virtual NVMe Device. (FAST'23).** Sang-Hoon Kim, Jaehoon Shim, Seongyeop Jeong, Jin-Soo Kim

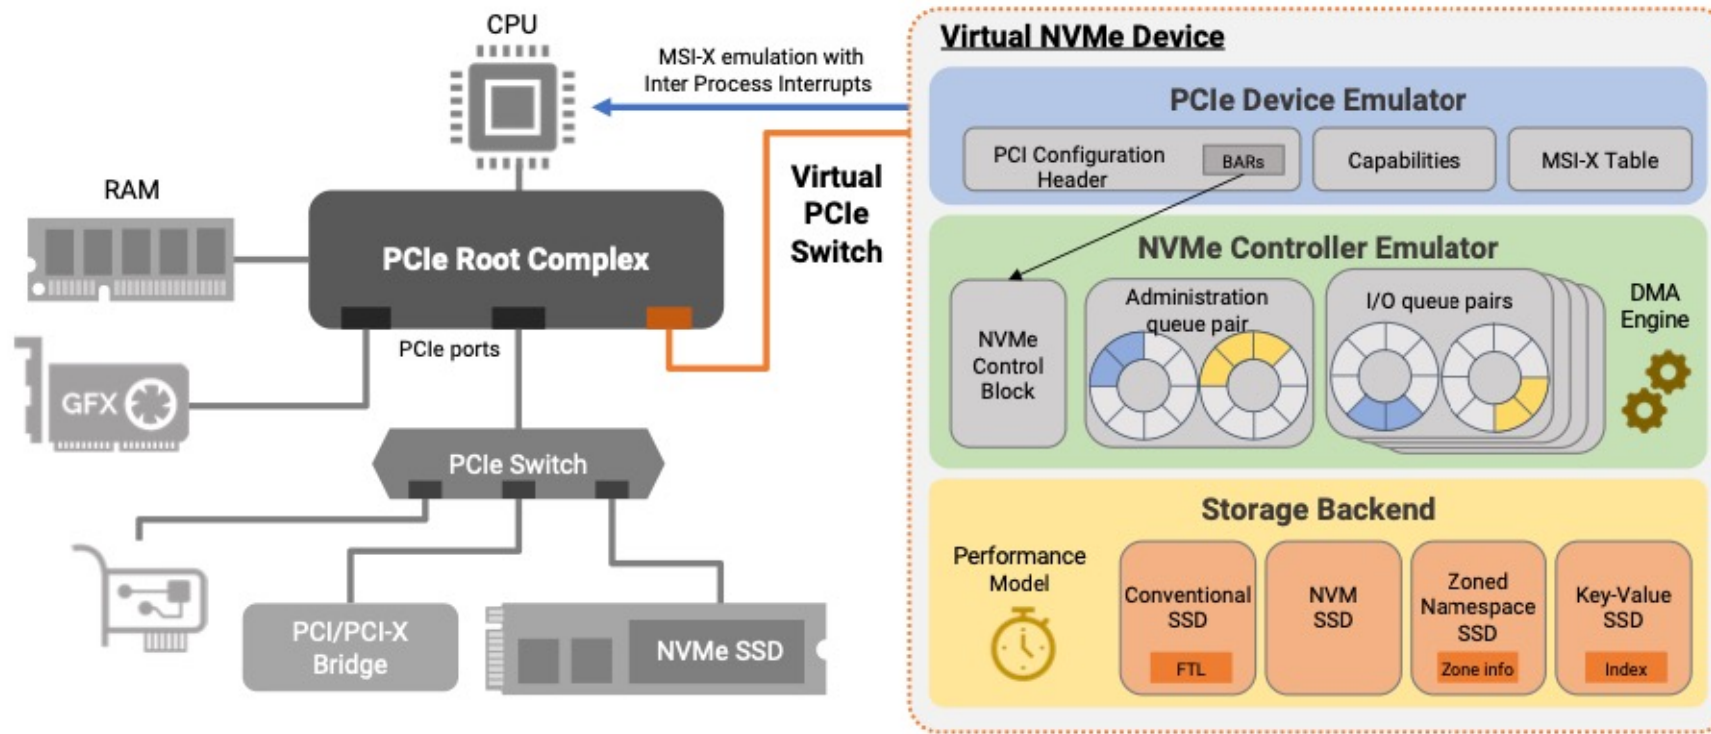


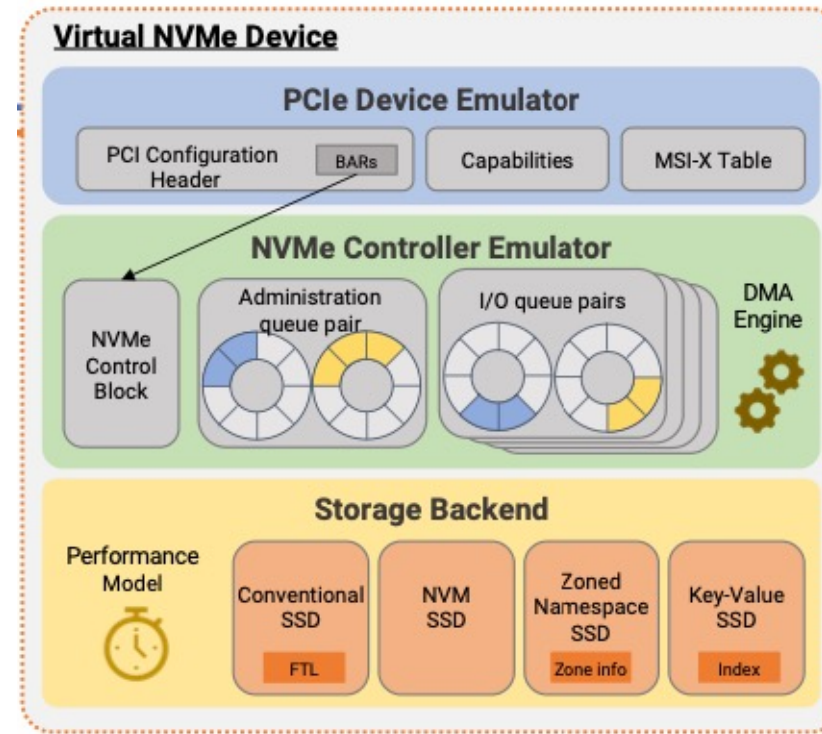
Figure 1: The overall architecture of NVMeVirt. NVMeVirt virtualizes a virtual NVMe device through the PCI bus and switch, so the device is seen as a native PCIe device from the host.

4. Newly related work in 2023

2. **NVMeVirt: A Versatile Software-defined Virtual NVMe Device. (FAST'23). Sang-Hoon Kim, Jaehoon Shim, Seongyeop Jeong, Jin-Soo Kim**

**FEMU NVMe
interface**

.VS



HDU – Center of Cloud computing

4. Newly related work in 2023

3. **DECC: Differential ECC for Read Performance Optimization on High-Density NAND Flash Memory.** (ASPDAC '23). Yunpeng Song, Yina Lv, Liang Shi

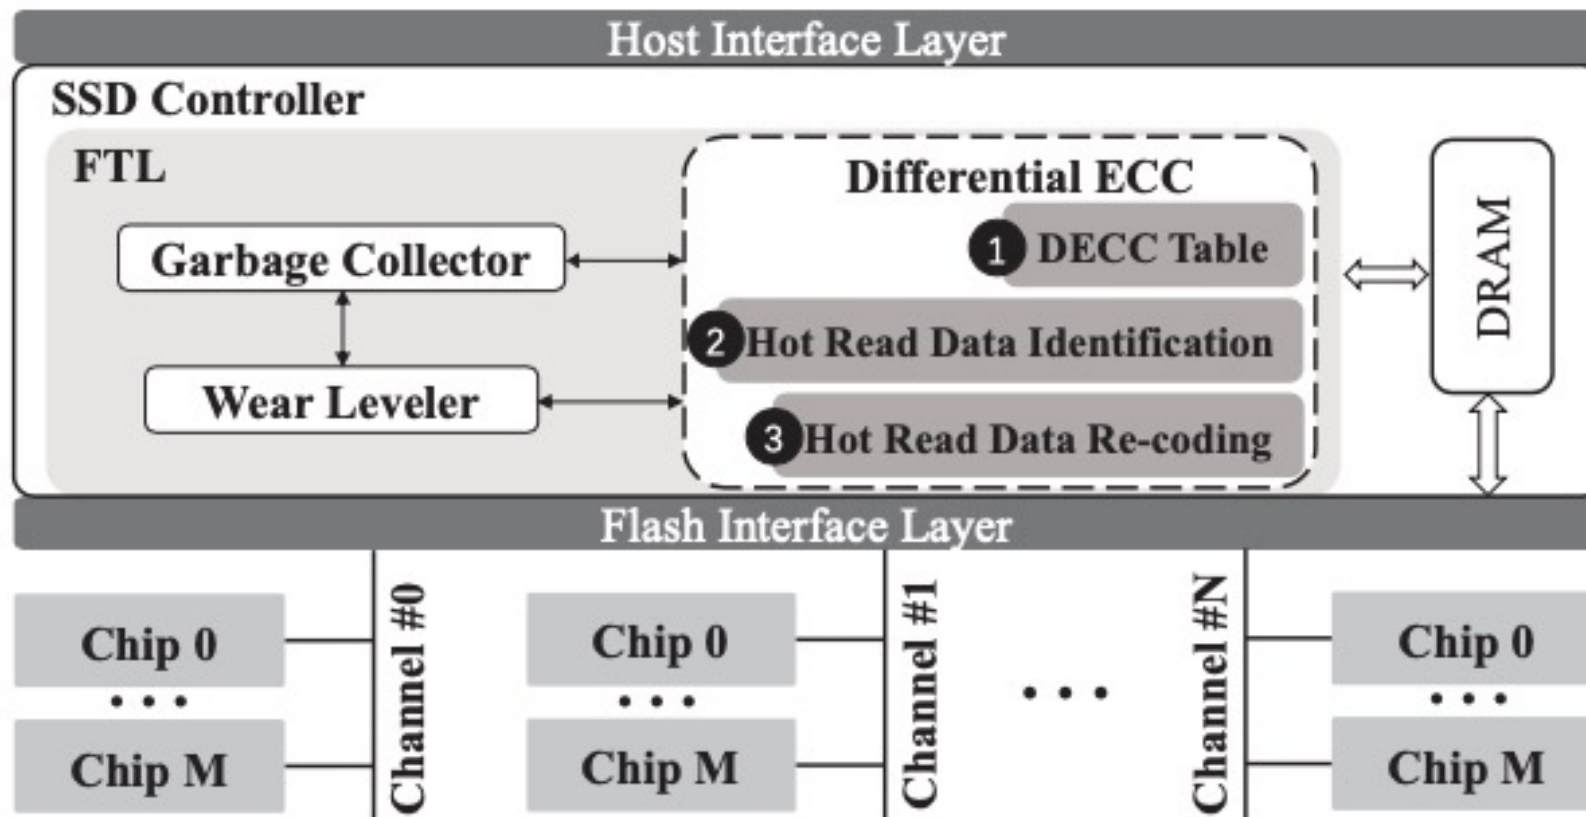


Figure 4: Overview of differential ECC.

4. Newly related work in 2023

3. **DECC: Differential ECC for Read Performance Optimization on High-Density NAND Flash Memory.** (ASPDAC '23). Yunpeng Song, Yina Lv, Liang Shi

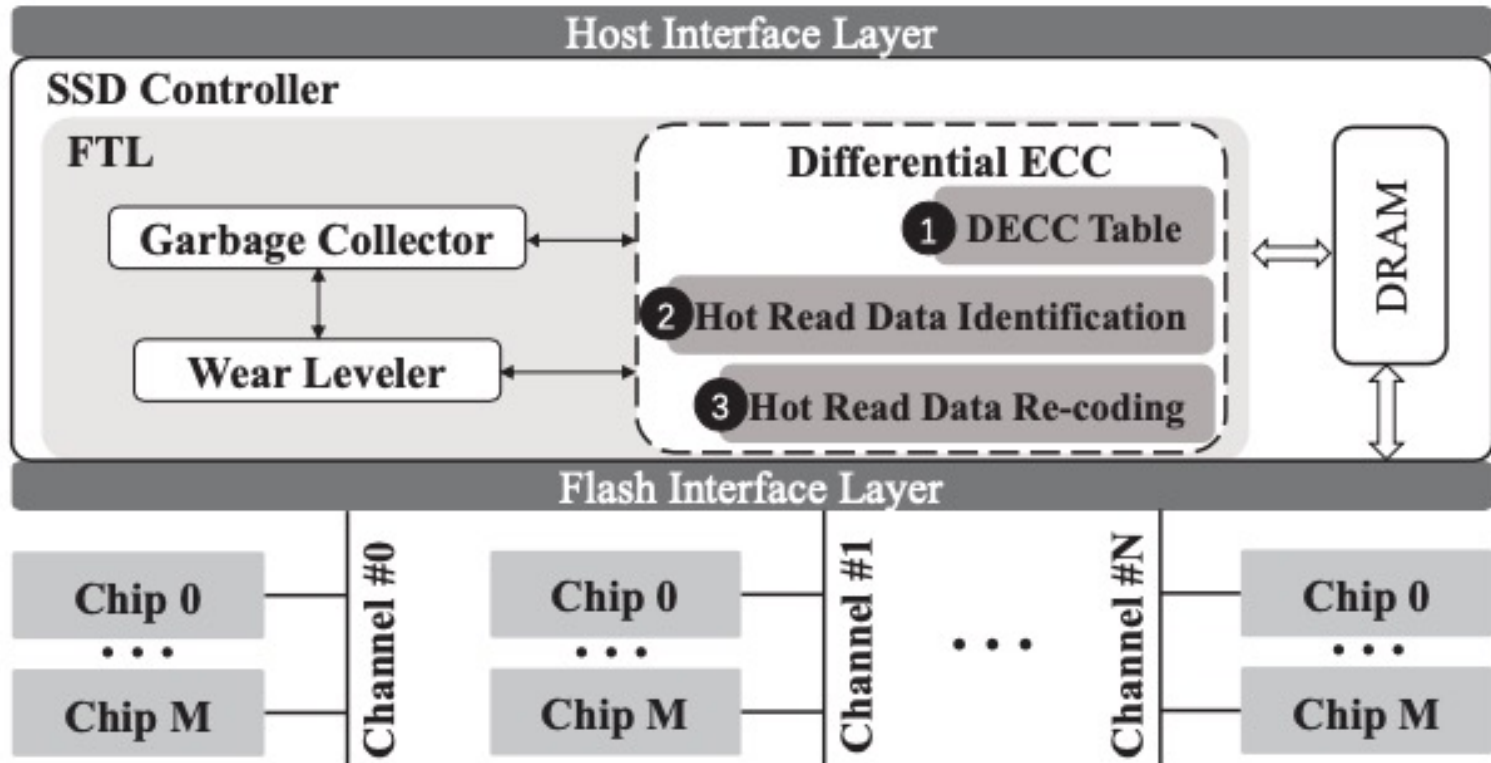


Figure 4: Overview of differential ECC.



FEMU
Kernel changes



FTL changes

HDU – Center of Cloud computing

4. Newly related work in 2023

4. **ConfZNS : A Novel Emulator for Exploring Design Space of ZNS. (SYSTOR'23).** Inho Song, Myounghoon Oh, Bryan S. Kim, Seehwan Yoo, Jaedong Lee, Jongmoo Choi

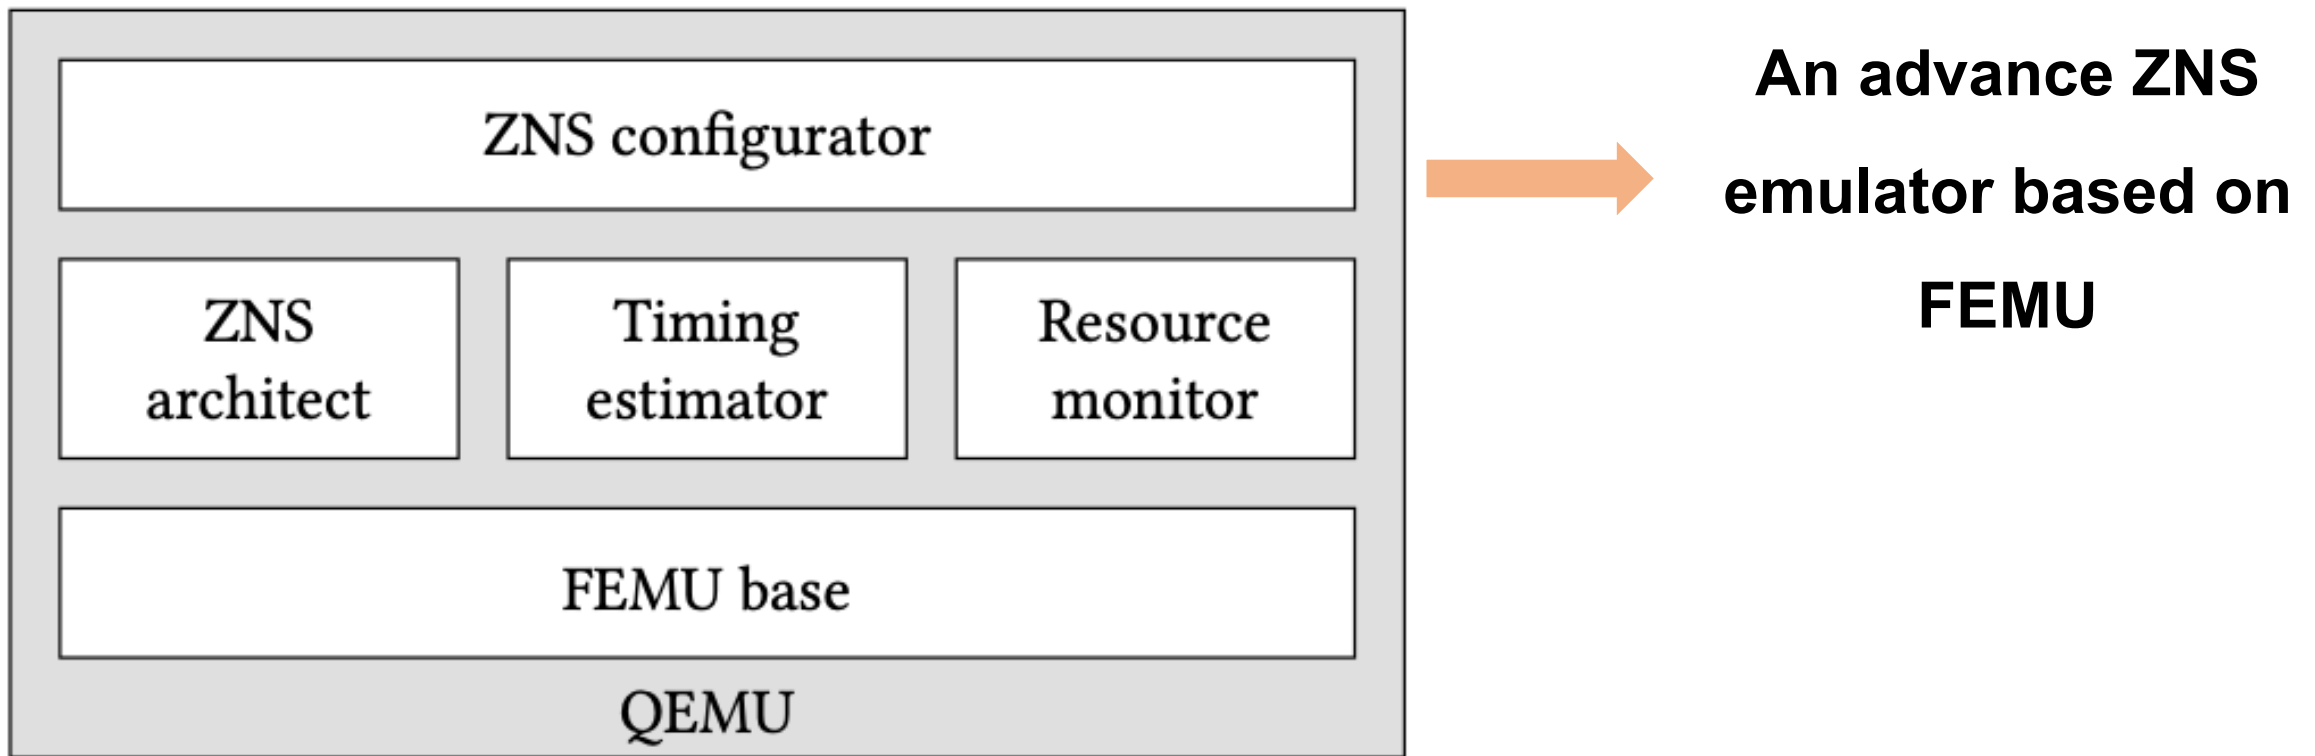


Figure 1: Key components of ConfZNS.

4. Newly related work in 2023

5. **Optimizing Data Migration for Garbage Collection in ZNS SSDs. (DATE'23).** Zhenhua Tan, Linbo Long, Renping Liu, Congming Gao, Yi Jiang, Yan Liu




Fig. 1. The method of in-storage data migration.

4. Newly related work in 2023

5. **Optimizing Data Migration for Garbage Collection in ZNS SSDs. (DATE'23).** Zhenhua Tan, Linbo Long, Renping Liu, Congming Gao, Yi Jiang, Yan Liu

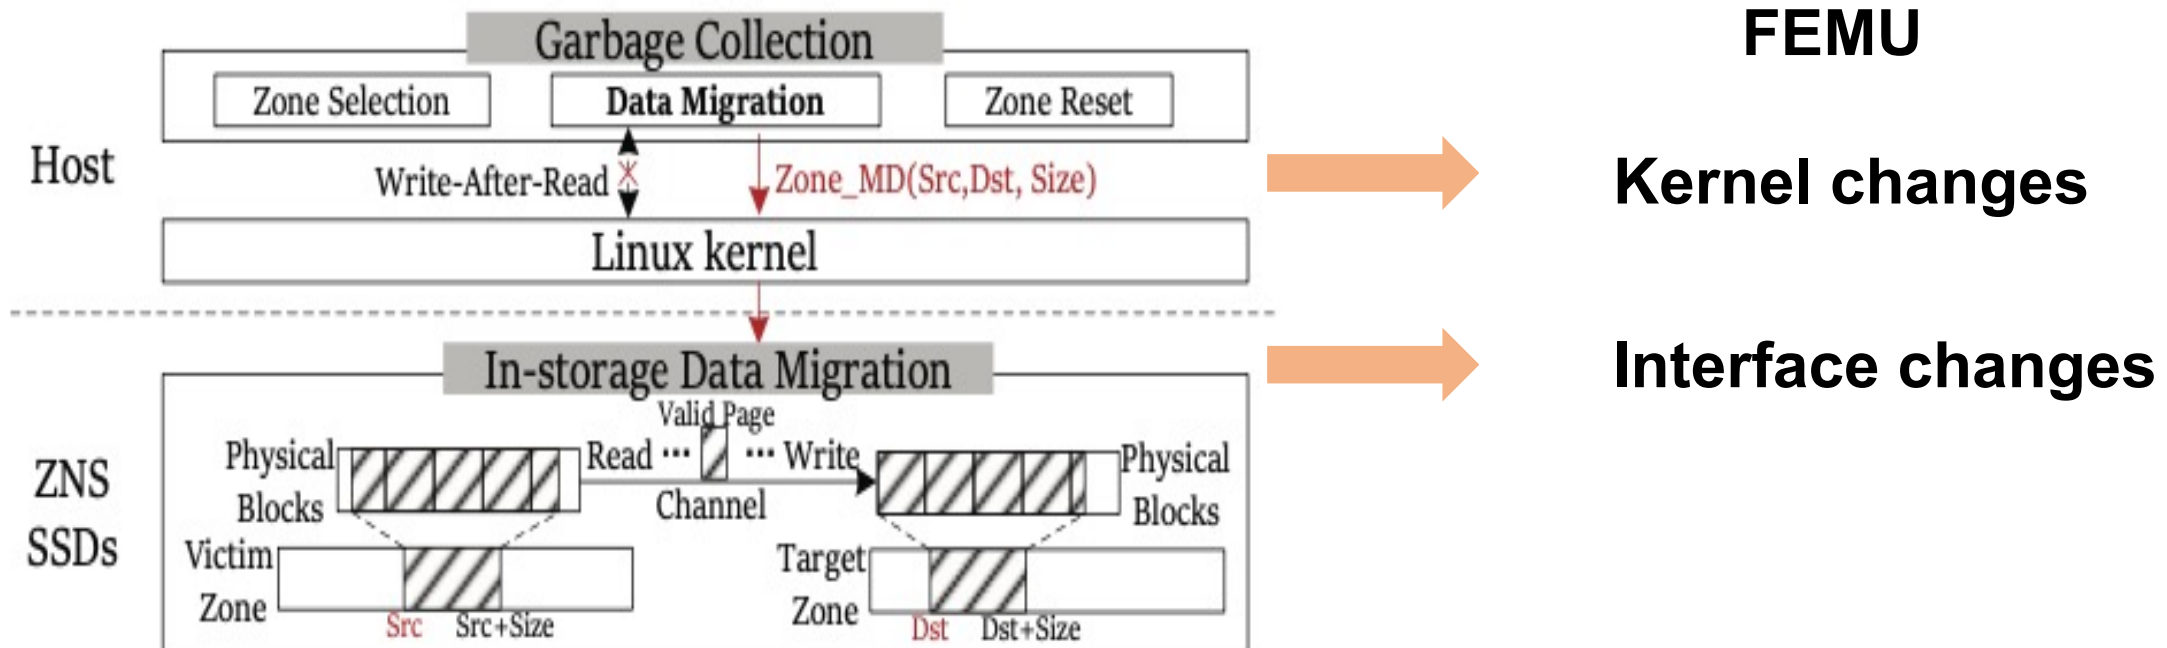


Fig. 1. The method of in-storage data migration.

4. Newly related work in 2023

6. Design of a High-Performance, High-Endurance Key-Value SSD for Large-Key Workloads. (IEEE CAL'23). Chanyoung Park, Chun-Yi Liu, Kyungtae Kang, Mahmut Kandemir, Wonil Choi




Fig. 4. Internal architecture of LK-SSD.




Fig. 2. Internal architecture of PinK.




Fig. 5. Relocating a set of values to the value log in LK-SSD+.

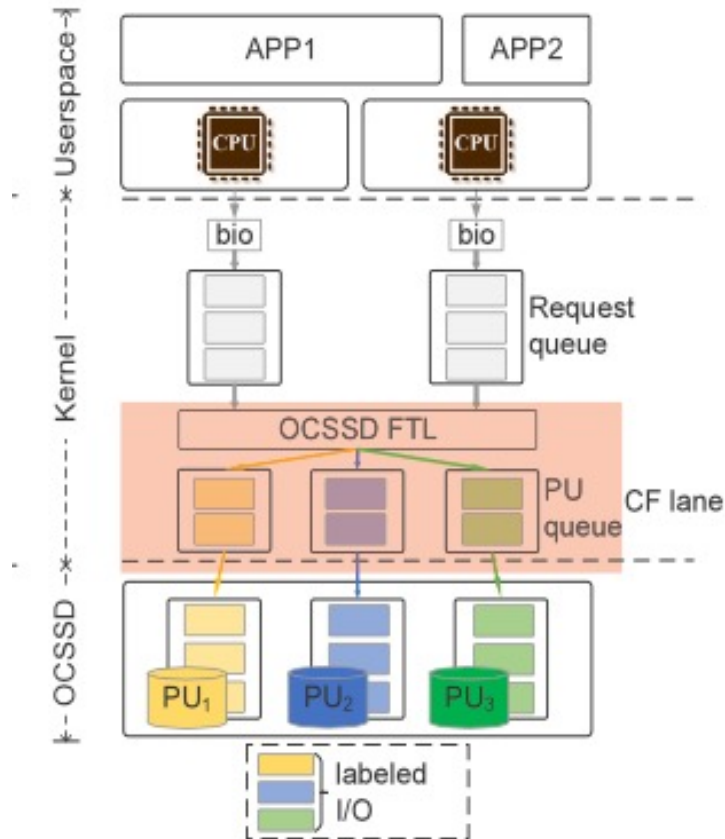
**Emulate three architecture on
FEMU,
Compare their performance**

HDU – Center of Cloud computing

4. Newly related work in 2023

7. CFIO: A conflict-free I/O mechanism to fully exploit internal parallelism for Open-Channel SSDs. (JSA'23).

Jinbin Zhu, Liang Wang, Limin Xiao, Lei Liu, Guangjun Qi



(b) The optimized I/O stack with CF lane.

FEMU

FTL changes:

*Optimized I/O stack with
CF line based on original
FEMU OC-SSD*

HDU – Center of Cloud computing

4. Newly related work in 2023

8. An Efficient F2FS GC Scheme for Improving I/O Latency of Foreground Application. (ICCE'23). Manjong Lee, Jonggyu Pa, Young Ik Eom

IV. EVALUATION

A. Experimental Setup

Table 1. Experiment Setup

Host Spec.	CPU: Intel(R) Xeon(R) CPU E5-2620 v4 @ 2.10GHz Memory: DDR4 64GB OS: Ubuntu 18.04.6 LTS
VM Spec.	Emulator: FEMU-v5.2 CPU: 4 SMP CPU Memory: 4GB OS: Ubuntu 20.04.4 LTS Kernel Version: Kernel v5.17 Storage: 12GB ZNS SSD

The detailed experiment setup is described in Table 1. To emulate ZNS storage, we use FEMU [15] with the default setting. To test various situations, we set various zone size from 16MB to 128MB. We use FIO [16] as the background application issuing write operations. The background application repeatedly creates/deletes 32KB ~ 64KB small files and 256KB ~ 1MB big files. For the foreground application, we use a program that issues 4KB write operations followed by fsync every 30ms.

***Emulate different zone size on
FEMU,
Compare their performance***

5. Future work

- 1. make FEMU running on my computer**
- 2. Read the Friendly Source Code of FEMU**
- 3. Find research direction**

Thank you for listening

Presented by Chen Haodong